Linux Advanced Systems Administration

Web Server - Virtual Host facility

- Our current Apache configuration supports a single web site, using the hostname as the URL
- In practise, most web servers would have to support many more than this, so how could this be achieved?:-
    * By having additional copies of *httpd*, each one looking after its own site.
    * By having a single copy of *httpd* that manages more than one site.
- The single copy option can be achieved using Apache's Virtual Hosts facility, and this is how the system should be configured.
- Virtual hosts can be IP-based or Name based (The preferred method)
- Imagine that we now have an additional new web site directory under */var/www* called *lintrain*.
- So we now have two sites:-
    * Our original */var/www/website/docs*, accessible simply as *http://hostname*
    * A new web site directory */var/www/lintrain/docs* with no server at present.
- Whar we will do is use the Virtual Hosts facility to serve each site, and use more conventional web server names.

Linux Advanced Systems Administration

Web Server - Virtual Host facility

- We need to place entries in the *httpd.conf* files for each site, as follows:- (Note that there are preferred places where these should be placed - have a look for the examples already in the file)

```
NameVirtualHost     192.168.200.100
<VirtualHost        www.fatrain.com>
ServerAdmin         root@carlton
DocumentRoot        /var/www/website/docs
ServerName          www.fatrain.com
ErrorLog            /var/www/logs/website_error_log
TransferLog         /var/www/logs/website_xfer_log
</VirtualHost>
<VirtualHost        www.lintrain.com>
ServerAdmin         root@carlton
DocumentRoot        /var/www/lintrain/docs
ServerName          www.lintrain.com
ErrorLog            /var/www/logs/lintrain_error_log
TransferLog         /var/www/logs/lintrain_xfer_log
</VirtualHost>
```

  - * The */var/www/logs* directory must be created.
  - * The *NameVirtualHost* directive specifies the IP address of the web server (In the example, this is *carlton* on 192.168.200.100)
  - * The hostnames *www.fatrain.com* and *www.lintrain.com* must be in */etc/hosts*

Linux Advanced Systems Administration

Web Server - Virtual Host facility

- Example hosts entry:-

192.168.200.100 carlton www.lintrain.com www.fatrain.com

- We must also add entries to *httpd.conf* to ensure that the sites are accessible:-

```
<Directory "/var/www/website">
   Options Indexes Includes FollowSymLinks
   AllowOverride None
   Order allow,deny
   Allow from all
</Directory>
<Directory "/var/www/lintrain">
   Options Indexes Includes FollowSymLinks
   AllowOverride None
   Order allow,deny
   Allow from all
</Directory>
```

  * Place these just after the VirtualHost entries just inserted.

- Note: Remove the original *DocumentRoot* directive.
- Sample sites are available under */net/duck/software/www*, in directories *website* (the existing site) and *lintrain*.
- Please use these to set up your own Apache virtual hosts, but use different hostnames to *www.fatrain.com* and *www.lintrain.com*. (PTO)

Linux Advanced Systems Administration

Web Server - Virtual Host facility - Notes

- As well as using the VirtualHost directive, we have also used a few other entries which have not been explained.
- We will be taking a look at these, including the use of CGI scripts, very soon, so please just accept their use for now!

Linux Advanced Systems Administration

Web Server - Controlling Virtual Hosts

- When running the web server with a number of web sites, httpd will start a number of child versions of itself to cater for the additional requests.
- There are a number of directives which can be used to set limits:-

MaxClients n

* Effectively limits the number of *httpd* servers which can be running. Default 150.

MaxRequestsPerChild n

* Each invocation of *httpd* can handle only this number of requests, then it is killed. This guards against software errors in *httpd* having too much effect. Default 100.

MaxSpareServers n
MinSpareServers n

* Keep at least MinSpareServers always active (default 5); do not allow more than MaxSpareServers to be waiting (DEFAULT 20)

StartServers n

* Start *StartServers* processes when *hpptd* is initially invoked, i.e. do not wait for demand to build up before starting them.

Linux Advanced Systems Administration

Web Server - CGI Scripts

- CGI (Common Gateway Interface) scripts are an important part of the webserver setup.
- CGI scripts, often written in *perl*, are run as a result of user interaction with the web site.
- This is an example HTML web page which activates a CGI script:-

```
<HTML>
<HEAD>
<TITLE>Training Example </TITLE>
</HEAD>
<BODY BGCOLOR=white text=blue link=green vlink=red>
<CENTER><H1>Training example </H1>
<FORM method=get action="http:/cgi-bin/demo.pl">
Please enter first data field <INPUT name=cat><BR>
Please enter second data field <INPUT name=dog><BR>
Then select the go button to see what your perl program does
<INPUT TYPE=submit Value="go!">
</FORM>
</BODY>
</HTML>
```

- The CGI script being called is *demo.pl*, located (apparently) under */cgi-bin*.

Linux Advanced Systems Administration

Web Server - CGI Scripts

- This is the CGI script:-

```
#!/usr/bin/perl
$_ = $ENV{"QUERY_STRING"};
($cat,$dog) = /^cat=(.*)&dog=(.*)$/;

print <<"PAGE" ;
Content-type: text/html

<HTML><BODY BGCOLOR=white text=black>
The strings the perl program received were:<BR>
First field: $cat<BR>Second field: $dog<P>
</BODY></HTML>
PAGE
```

- The location of the HTML page is, of course, within a directory under the web site.
- CGI scripts can also be held within the web site directory, but this could cause security problems if the script sources are accessible from outside.
- Ideally, they should be held elsewhere on the web server, in a directory that cannot be accessed as part of the web site.
- In our previous (initial) example, we defined the location of this directory using the *ScriptAlias* directive.

Linux Advanced Systems Administration

Web Server - CGI Scripts - Location

- Here is how we defined it:-

ScriptAlias /cgi-bin/ "**/var/www/website/cgi-bin/**"

* Then, just below:-

<Directory "**/var/www/website/cgi-bin**">
AllowOverride None
Options ExecCGI
Order allow,deny
Allow from all
</Directory>

- When a web link to a CGI script specifies */cgi-bin* then the web server can locate it under the specified directory.
- To test this facility, we need to add our CGI script and web page to our existing setup - we will do this using the *lintrain* web site.

Linux Advanced Systems Administration

Web Server - CGI Scripts - Example

- To begin with, copy the script and the web page from *duck*:-

```
# cd   /net/duck/software/www
# ls
CGI-example   lintrain   website
# cp   -r   CGI-example   /var/www
# cd   /var/www/CGI-example
# cp   profile.html   ../lintrain/docs
# cp   demo.pl   ../cgi-bin
```

  * Remember to make the script execuatble!

```
# chmod  +x   ../cgi-bin/demo.pl
```

- To keep things consistent, and avoid having multiple *cgi-bin* directories, we will move all the *website/cgi-bin* stuff into the *cgi-bin* in */var/www*:-

```
# cd  /var/www/website/cgi-bin
# mv   *   /var/www/cgi-bin
```

- Now edit the *index.html* under the *lintrain/docs* directory to add a link to the new web page:-

```
# cd  /var/www/lintrain/docs
# vi  index.html
```

  * Add the following in some appropriate position:-

```
</P>
<P><B>
Click <A href="profile.html"><STRONG> HERE </STRONG></A> to see the cgi
example form.  </B>
</P>
```

Linux Advanced Systems Administration

Web Server - CGI Scripts

- Now we must edit *httpd.conf*:-

ScriptAlias /cgi-bin/ "**/var/www/website/cgi-bin/**"

* Change to:-

ScriptAlias /cgi-bin/ "**/var/www/cgi-bin**”

* Then, just below:-

```
<Directory "/var/www/website/cgi-bin">
   AllowOverride None
   Options ExecCGI
   Order allow,deny
   Allow from all
</Directory>
```

* Change to:-

```
<Directory "/var/www/cgi-bin">
   AllowOverride None
   Options ExecCGI
   Order allow,deny
   Allow from all
</Directory>
```

- Now stop and restart *httpd*:-

/etc/rc.d/init.d/httpd stop

/etc/rc.d/init.d/httpd start

- Now point your browser at *www.lintrain.com* once again:-

Linux Advanced Systems Administration

Web Server - CGI Scripts

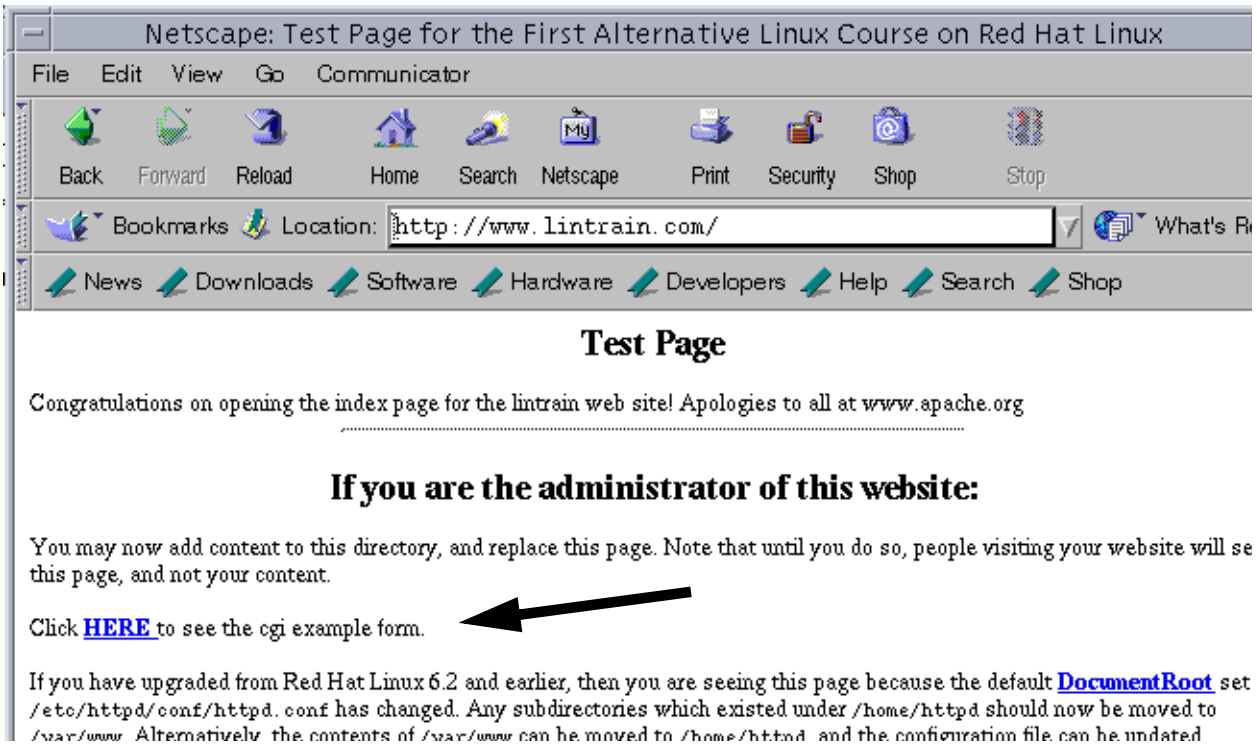

- The newly-inserted link is indicated.
- Click on the link to see:-

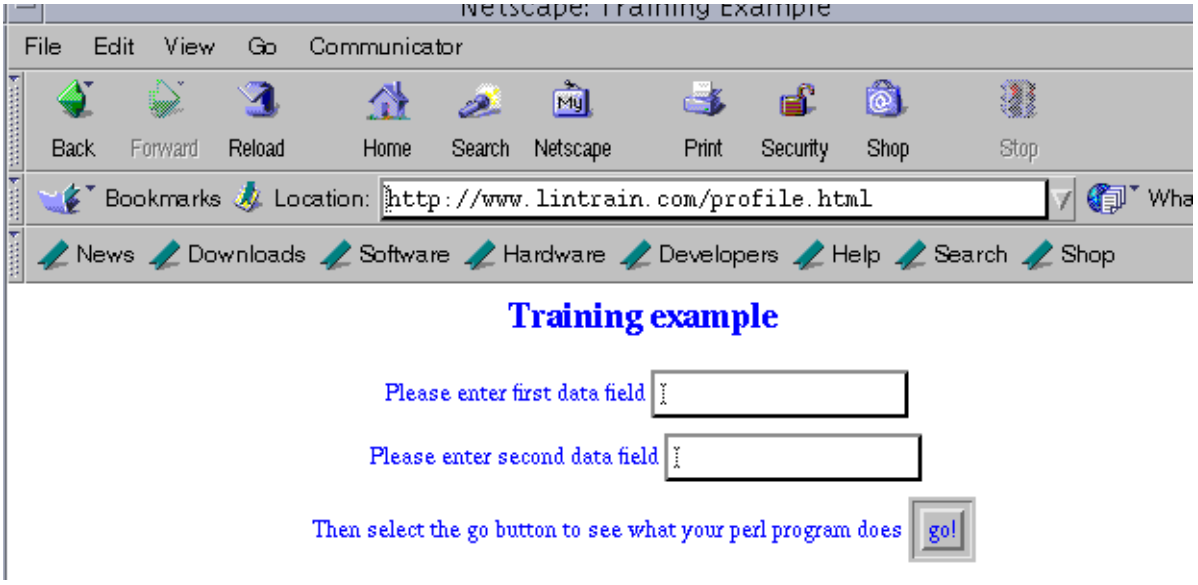

- Complete the fields and click on Go.

Linux Advanced Systems Administration

Web Server - CGI Scripts

- I entered *Harwell* and *Didcot*, so I will see:-

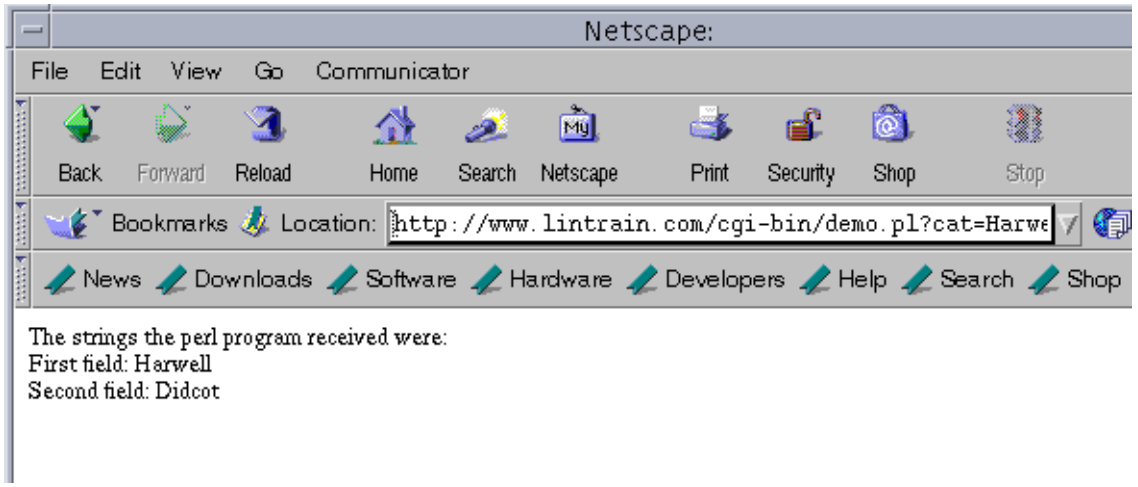

- So we have now integrated the CGI scripting facility into our web server, and can now confidently add more scripts to make our site much more powerful.
- There is a lot to CGI script writing, which (sorry!) is beyond the scope of this course.